

```
1
2
3
4 Breedify{ Dog breeds
5 classification
6
7
8     [with Machine Learning
9     and Deep Learning]
10
11
12 }
13
14
```

Introduction

Dogs are one of the most successful and diverse species that populate our planet, in fact, there exist between 195 and 500 dog breeds around the world.

Many of these breeds are totally **different** from one another, while many other present remarkable **similarity**, and are considered to be different just for some subtle details.

For these reasons, the classification of dogs into different breeds is a **very challenging task**.

In this project we tried two different approaches based on Machine Learning, and Deep Learning respectively, to understand which was the most effective to tackle this challenge.



Table Of Contents {

01 Data Collection

02 Machine Learning Model with Feature Extraction

< HOG vs LBP feature extraction>

< SVM, Random Forest, AdaBoost>

03 Deep Learning Model with CNNs

< MobileNet CNN>


01 { Data Collection

Our data was collected from a publicly available dataset that we found for free on Kaggle.

Specifically, we used the Stanford Dogs Dataset, which contains 20,580 images of 120 dogs from around the world.

Stanford Dogs Dataset

Over 20,000 images of 120 dog breeds





Machine Learning with Feature Descriptors{

< In order to build our Machine Learning model, we used the **scikit-learn** library.

We shuffled our data, and split them in training (60%), validation (10%), and test (30%) sets>

```
[ ] np.random.shuffle(images)
```

```
[ ] trainset = images[:int(0.6*len(images))]  
    valset = images[int(0.6*len(images)):int(0.7*len(images))]  
    testset = images[int(0.7*len(images)):]
```

```
print('Total: {} splitted in Train: {}, Val: {} and Test: {}'.format(len(images), len(trainset), len(valset), len(testset)))
```

```
Total: 20580 splitted in Train: 12348, Val: 2057 and Test: 6175
```

Feature Descriptors{

We used 2 visual Feature Descriptors: **HOG** and **LBP**, and defined the `extract_features()` function to perform feature extraction for each image:

```
def extract_features(images, feat_type, img_size):

    labels = []
    features = []

    for image in images:

        # open the image
        img = cv2.imread(image, 0)

        # resize the image
        img = cv2.resize(img, (img_size, img_size))

        # compute the features
        if feat_type == 'hog':
            feat = hog(img, orientations=8, pixels_per_cell=(4, 4), cells_per_block=(1, 1))
        elif feat_type == 'lbp':
            feat = np.ravel(local_binary_pattern(img, P=100, R=5))
        else:
            raise NotImplementedError('Not implemented feature!')

        # append features and labels
        features.append(feat)

        labels.append(get_labels(image))

    return features, labels
```

ML Classifiers{

For the classification task, we used 3 different classifiers:
Support Vector Machines, **Random Forest**, and **AdaBoost**.

We defined the **decide(clf)** function, which selects one classifier at a time and runs the classification with different features and image size combinations:

```
def decide_clf (clf):  
    if clf == "SVM":  
        clf = svm.SVC(gamma=0.001, C=100., kernel='linear', class_weight='balanced', verbose=False, probability=False)  
    elif clf == "RF":  
        clf = RandomForestClassifier(n_estimators=100, criterion='gini', max_depth=None, min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0,  
    elif clf == "Adaboost":  
        clf = AdaBoostClassifier(estimator=None, n_estimators=50, learning_rate=1.0, algorithm='SAMME.R', random_state=None)  
    return clf
```

1 ML Classifiers { e {}'.format(e1, feature_type, img_size, accuracy_score(test_y, y_pred), t2-t1))

2

```
sizes = [32, 64, 128]
features = ['hog', 'lbp']
types_clf = ["SVM", "RF", "Adaboost"]

def decide_clf (clf):
    if clf == "SVM":
        clf = svm.SVC(gamma=0.001, C=100., kernel='linear', class_weight='balanced', verbose=False, probability=False)
    elif clf == "RF":
        clf = RandomForestClassifier(n_estimators=100, criterion='gini', max_depth=None, min_samples_split=2, min_samples_leaf=1, min_weight_fraction=0.01)
    elif clf == "Adaboost":
        clf = AdaBoostClassifier(estimator=None, n_estimators=50, learning_rate=1.0, algorithm='SAMME.R', random_state=None)
    return clf

for img_size in sizes:
    for feature_type in features:
        t1 = time.time()
        train_x, train_y = extract_features(trainset, feature_type, img_size)
        val_x, val_y = extract_features(valset, feature_type, img_size)
        test_x, test_y = extract_features(testset, feature_type, img_size)
        for el in types_clf:
            clf = decide_clf(el)
            clf.fit(train_x, train_y)
            clf.score(val_x, val_y)
            y_pred = clf.predict(test_x) #to predict the class of the given data.
            t2 = time.time()
            print('Final Accuracy for {} classifier with {} feature descriptor and {} image size: {:.3f}. Elapsed time {}'.format(el, feature_type, img_size, accuracy_score(test_y, y_pred), t2-t1))
```

feature extraction

training, validation, and prediction

SVM Accuracy Results {

| Feature | Image size | Accuracy | Elapsed time |
|------------|----------------|--------------|----------------|
| HOG | 32x32 | 0.021 | 153.256 |
| LBP | 32x32 | 0.016 | 450.756 |
| HOG | 64x64 | 0.032 | 242.656 |
| LBP | 64x64 | 0.016 | 1564.398 |
| HOG | 128X128 | 0.041 | 690.136 |
| LBP | 128x128 | 0.018 | 7812.561 |

Random Forest Accuracy Results {

| Feature | Image size | Accuracy | Elapsed time |
|---------|------------|----------|--------------|
| HOG | 32x32 | 0.020 | 212.769 |
| LBP | 32x32 | 0.021 | 712.115 |
| HOG | 64x64 | 0.018 | 543.993 |
| LBP | 64x64 | 0.015 | 2100.234 |
| HOG | 128x128 | 0.018 | 1802.291 |
| LBP | 128x128 | 0.016 | 9012.543 |

AdaBoost Accuracy Results {

| Feature | Image size | Accuracy | Elapsed time |
|---------|------------|----------|--------------|
| HOG | 32x32 | 0.012 | 253.476 |
| LBP | 32x32 | 0.014 | 923.209 |
| HOG | 64x64 | 0.015 | 670.137 |
| LBP | 64x64 | 0.011 | 2910.388 |
| HOG | 128x128 | 0.015 | 2456.893 |
| LBP | 128x128 | 0.011 | 12546.854 |

Deep Learning with Convolutional Neural Networks{

For our Deep Learning model, we used the **TensorFlow** library with the Keras interface.

To define the architecture of our Convolutional Neural Network, we used the **MobileNet CNN**, adapted to our case.

```
# Our model
model = tf.keras.applications.MobileNet(
    input_shape=None,
    alpha=1.0,
    depth_multiplier=1,
    dropout=0.001,
    include_top=True,
    weights="imagenet",
    input_tensor=None,
    pooling=None,
    classes=1000,
    classifier_activation="softmax"
)
```

```
# here we create the new Dense layer where input is the output of the second last layer
output = Dense(120, activation='softmax', name='predictions')(model.layers[-2].output)

# Then create the corresponding model
model = Model(model.input, output)
```

1 CNN code{ 2

```
epochs = 30

callbacks = [
    keras.callbacks.ModelCheckpoint("save_at_{epoch}.h30"),
    tf.keras.callbacks.TensorBoard(log_dir="logs", write_graph=True, write_images=False, update_freq="epoch",)
]

model.compile(
    optimizer=keras.optimizers.SGD(1e-3),
    loss="categorical_crossentropy",
    metrics=["accuracy"],
)
```

10

11 Here we define the epochs, callbacks, optimizers and loss function.

12


The final results were: 63s 191ms/step - **loss: 0.3028** - **accuracy: 0.9307**

13

14

CNN Results{

Final Accuracy: 0.9307



CNN Results{

Loss: 0.3028

